

Card Counting Blackjack: Phase 2

Michalis Lamprakis 2020030077
Dimitris Ilia 2020030200

Project context

The goal of the second phase of the project is to make a more realistic Blackjack environment and hopefully "beat" the casino. The environment now includes betting option, double down action and 3/2 payout in natural blackjack. The project is structured in 2 subtasks: i) fixed 1\$ bet at the start of each round and ii) agent to chooses the initial bet.

Implementantion

The whole implementantion is based on **BlackjackEnv** class that we made from scratch ,with the help of online tools and resources, and contains all the necessary methods to simulate the game. The provided code contains detailed comments.

Task 1

In this task, there is a fixed bet of 1\$ at the start of each round with the option to double down. Natural blackjack pays 1.5 times the bet. We train an agent to find the optimal policy (with and without Card Counting) using tabular Q-learning and Bellman equation.

$$Q(s, a) \leftarrow Q(s, a) + \alpha \left(r + \gamma \max_{a'} Q(s', a') - Q(s, a) \right)$$

The state space consists of

- (player's total, dealer's upcard, usable ace, can_double) for the basic agent
- (player's total, dealer's upcard, usable ace, can_double, hi_lo_bucket) for the hi-lo environment

can_double is a boolean indicating whether the player can double down.

After experimenting with different parameters, we decided to use

- **Alpha**: decayed linearly from 0.1 to 0.01 across 200,000 episodes
- **Gamma**: 1.0
- **Epsilon**: decayed linearly from 1 to 0.05 across 1M episodes

The agent is trained for 5M episodes (for both cases) and evaluated for 100K episodes with the learned policy, we capture the following results:

	Basic Agent	Hi-Lo Agent
Average profit per game	-0.02\$	-0.01\$
Average bet size	1.00\$	1.00\$
Double-down rate	1.36%	2.37%
Double-down win rate	64.75%	59.53%
Win	43.25%	43.74%
Draw	9.13%	8.72%
Lose	47.62%	47.55%

We observe that the hi-lo agent performs better than the basic agent, which is expected as it uses card counting to adjust its strategy, but still both agents have a negative average profit per game. Also the hi-lo is more confident to double down, as it has more information about the remaining cards.




Figure 1: Average reward per episode for the Basic Agent (left) and Hi-Lo Agent (right).

To monitor the learning progress of our Q-learning agent, we tracked the average reward per episode throughout training. At the beginning of training, the agent explores heavily, resulting in highly variable returns. As training progresses and the exploration rate decays, the agent converges to a more stable policy.

We also did the comparison of our learned policy with the proven best strategy for Blackjack and we observed 96-97% compliance which is a proof that our agent has learned the optimal policy.

Execution details

For this task there is a menu with the following options:

- **1:** Play manually against the dealer
- **2:** Train the basic Q-agent
- **3:** Evaluate the basic agent
- **4:** Compare the learned policy with the proven best strategy
- **5:** Train the hi-lo agent
- **6:** Evaluate the hi-lo agent
- **7:** Exit

Task 2

For the second task, we can now choose the bet size at the start of each round. We implemented two separate neural networks with different roles and inputs: a Play Network (DQN) responsible for decision-making during gameplay, and a Bet Network (BetNet) responsible for determining optimal bet sizes based on card count.

The **Play Network** is a Deep Q-Network (DQN) that approximates the optimal action-value function $Q(s, a)$ for the Blackjack play policy. It learns to choose between Hit, Stick, or Double Down, based on the current game state. It is a fully connected neural network with 2 hidden layers, the first one contains 128 neurons and the second 64. Both, use ReLU activation.

- **Input:** player's total, dealer's upcard, usable ace, can_double, true count (state vector of size 5)
- **Output:** Q-values for HIT (0), STICK (1), DOUBLE (2)

We have chosen the following hyperparameters for the DQN:

- **Alpha:** 1e-4
- **Gamma:** 1.0
- **Epsilon:** decayed linearly from 1.0 to 0.05 across 1M episodes
- **Batch Size:** 128
- **Replay Buffer Size:** 200,000 transitions
- **Target Network Update Frequency:** every 5,000 steps

The **BetNet** is a separate neural network that learns a bet sizing policy based only on the true count. It outputs a Distribution over 10 discrete bet sizes (\$1–\$10), allowing the agent to bet more when the deck is favorable and less when it is not. It is also a smaller fully connected neural network with 2 hidden layers, the first one contains 64 neurons and the second 32. The activation function is ReLU for both layers.

- **Input:** true count (state vector of size 1)
- **Output:** Distribution over bet sizes (10 outputs)

And the hyperparameters for the BetNet:

- **Alpha:** 1e-4
- **Gamma:** 1.0
- **Epsilon:** decayed linearly from 1.0 to 0.05 across 1M episodes
- **Batch Size:** 128
- **Replay Buffer Size:** 200,000 transitions
- **Target Network Update Frequency:** every 5,000 steps

After training both networks for 1.6M episodes, we evaluate the agent for another 100K episodes and we capture the following results (for 1 deck and 4 decks environment):

	1 Deck	4 Decks
Average profit per game	0.09\$	0.04\$
Average bet size	5.13\$	4.15\$
Expected return per unit bet	0.0182	0.0104
Double-down rate	8.79%	6.63%
Double-down win rate	55.73%	54.95%
Win	44.00%	43.48%
Draw	8.27%	8.81%
Lose	47.73%	47.71%

We observe that the agent performs significantly better than the previous task, achieving a positive average profit per game, which means the agent is statistically profitable and effectively beating the casino under these settings. The agent is more confident to double down, as it has more information about the remaining cards and it can adjust its bet size according to them. Without the betting network, the DQN policy would be almost the same as Task1. So the key innovation here is the BetNet, which allows the agent to adjust its bet size based on the true count. This means the agent bets more when the deck is favorable, just like a real card counter. We can see exactly this in **Figure 2**. Agent maximizes its bet size when the true count is high, and minimizes it when the true count is low. This is a clear indication that the agent has learned to use card counting effectively to adjust its betting strategy.

	€1	€2	€3	€4	€5	€6	€7	€8	€9	€10
-10	2260.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0
-9	593.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
-8	1301.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
-7	1032.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
-6	2345.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
-5	2286.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
-4	4184.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
-3	984.0	4275.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
-2	0.0	8928.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
-1	1040.0	12106.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
+0	2982.0	1479.0	0.0	0.0	14731.0	0.0	0.0	0.0	0.0	0.0
+1	0.0	8893.0	0.0	0.0	0.0	0.0	0.0	0.0	3748.0	0.0
+2	0.0	3204.0	0.0	2024.0	0.0	0.0	0.0	557.0	0.0	2463.0
+3	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	4746.0
+4	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	4083.0
+5	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	2252.0
+6	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	2268.0
+7	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1101.0
+8	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1320.0
+9	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	561.0
+10	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	2253.0

Figure 2: BetNet's distribution over bet sizes for different true counts.

The following plot (left) shows the moving average reward per episode during the training of the DQN play policy. Initially, the agent performs poorly. As training progresses, the average reward steadily improves and stabilizes, reflecting the agent's convergence toward a near-optimal policy. The right plot shows the TD-loss over time that steadily decreases during training, indicating improved accuracy in the agent's Q-value predictions.




Figure 3: Average reward for the DQN agent (left) and the TD-loss over time (right).

Execution details

Menu options:

- **1:** Train DQN agent with BetNet
- **2:** Evaluate
- **3:** Compute confidence interval
- **4:** load NNs and evaluate (in case we save our models)
- **5:** Exit

All the above are implemented in **PHASE_2.ipynb** file.

Conclusion

With this project we have successfully implemented a realistic Blackjack environment and trained agents that can play the game effectively. Started in Phase 1 with a basic agent achieving a negative average reward of **-0.06\$** per game and ended up in Phase 2 with an "advantage" agent that can adjust its betting strategy and achieve a positive average reward up to **0.09\$** per game which is pretty impressive.

Notes

In phase 1 while at first we used a learning rate of 0.1, as you told us and then we found out ourselves, it was too high and it caused instability in the training process, so we decided to use 0.01. As we can see in the following plots, the average reward stabilizes and also we tracked agent's chosen action evolved for three high-impact states:

- (18,10,0)
- (13,5,0)
- (15,10,0)

The agent initially fluctuated but converged.

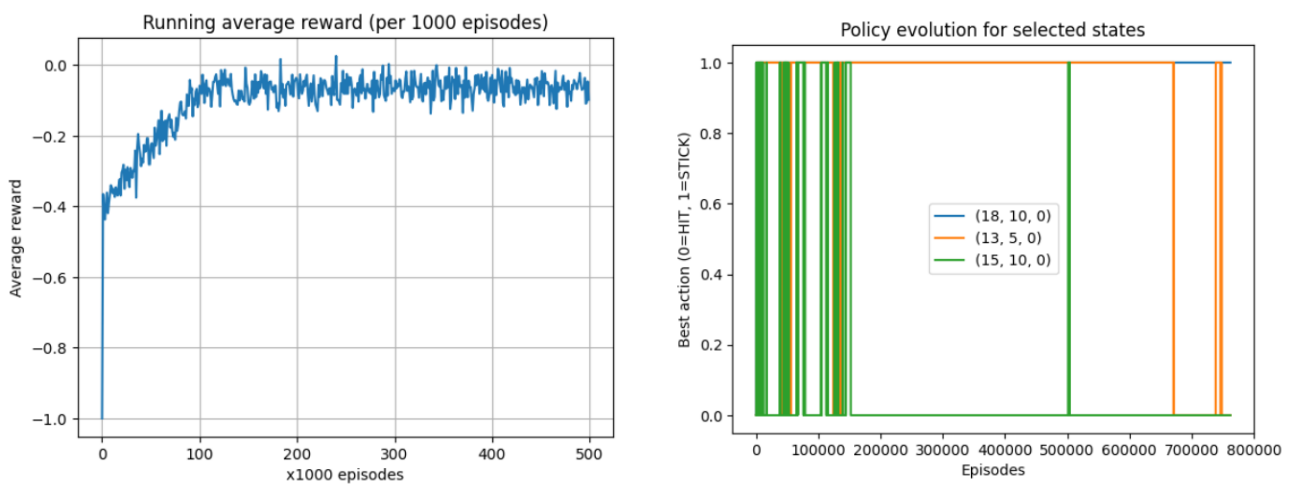


Figure 4: Average reward per episode for the Basic Agent with learning rate 0.01 (left) and the agent's chosen actions for three high-impact states (right).